

**Lab 3 – Address Translation + Memory Mapped Files**  
Operating Systems SFWRENG 3SH3, Winter 2025  
**Prof. Neerja Mhaskar**

1. **You must show your working solution of this lab to the TA for a grade.**
2. For Mac M1, M2, and M3 users (all Macs with 64-bit ARM CPUs), you need to install UTM for virtualization: <https://mac.getutm.app/>
3. The TA will check your solution and will quiz you on your work. After which they will enter your mark and feedback on Avenue.
4. If you do not show your work to your Lab TA, you will get a zero (unless you provide an MSAF, in which case this lab's weight will be moved to Assignment 2).
5. It is your responsibility to connect with your Lab TA to get a grade and ensure that your grade has indeed been posted on Avenue.

**Outline: PART I**

Assume that a system has a 32-bit virtual address with a 4-KB ( $=2^{12}$ ) page size. The physical memory address is also a 32-bit address. Consider a small program that needs only 8 pages of memory. Below is the page table for this program.

**Page Table:**

Frame Number
6
4
3
7
0
1
2
5

In these labs you are to write a C (labAT.c) program that simulates an MMU's (memory management unit's) address translation capability. To simulate a program's memory address requests, we use the text file named **labaddr.txt**. This file can be downloaded from the course website under Content->Labs->Lab 3 material. This file contains a sample of **20** logical addresses generated for this program. You are to read these addresses and output the following for every address:

Logical/virtual address, its corresponding page number and offset, and its corresponding physical address.

**There are three parts to this lab:**

1. Reading from a file.
2. Given a logical/virtual address output its page number and offset
3. Given the page table and the logical address output its corresponding physical address.

It is recommended that you approach these labs in the above order; that is, first read all the logical addresses from the file and simply output it to the terminal. Then compute the page number and offset and output these details for each logical address in the labaddr.txt file on the terminal. Finally, compute the physical address and display it on the terminal.

**Part I.1 – Reading from a file.**

1. Use the C library function `openf()` to open the '**labaddr.txt**' file. Since you will be simply reading from this file choose the 'r' (read) option.

Eg: `FILE *fptr = fopen("labaddr.txt", "r");`

2. To read a logical address from the file use the `fgets()` function. Since the logical addresses are no more than 10 characters long, you can read and store just 10 characters at a time. Sample code is:

```
#define BUFFER_SIZE 10;
char buff[BUFFER_SIZE];
//Read from labaddr.txt till you read end of file.
while(fgets(buff, BUFFER_SIZE, fptr) != NULL){...}
```

3. After reading a logical address print it to the terminal.
4. It is important that you close the file after you are done reading all the logical addresses from the file.

Sample code: `fclose(fptr);`

**Part I.2 – Given a logical address compute the page number and offset.**

1. Define `PAGE_NUMBER_MASK`, `OFFSET_MASK` etc. as macro definitions. Sample code is below, where you need to fill in the appropriate values in the blanks in your program.

```
#define OFFSET_MASK _____
#define PAGES _____
#define OFFSET_BITS _____
#define PAGE_SIZE _____
```

2. For each logical address compute, the page number and offset using **bitwise operators in C** (See notes at the end of Part I).

3. Print the logical address and its corresponding page number and offset to the terminal.

**Part I.3 – Given the logical address and page table compute the corresponding physical address.**

1. Define the page table as an integer array and store all the frame numbers as shown in the page table under the outline section of this document.

Eg: `int page_table[PAGES] = {6,4,3,7,0,1,2,5};`

2. After computing the page number (p) and offset (o), extract the frame number for the page (p) from the page table.

3. Using the frame number compute the corresponding physical address **using bitwise operators in C**.

3. Print the physical address along with the logical address and its corresponding page number and offset from PART II to the console.

4. Compile your program without errors and show the program's output to your TA.

**Important:**

1. Note that for your program to run correctly (and to avoid segmentation faults) it is important that you use correct data types for page number, frame\_number, virtual and physical addresses and offset.
2. Refer to lecture notes on paging. This content is explained under the lecture notes slides on Main memory (Chapter 9).

Correct Program output: ./labAT

Virtual addr is 19986: Page# = 4 & Offset = 3602. Physical addr = 3602.  
Virtual addr is 16916: Page# = 4 & Offset = 532. Physical addr = 532.  
Virtual addr is 24493: Page# = 5 & Offset = 4013. Physical addr = 8109.  
Virtual addr is 8198: Page# = 2 & Offset = 6. Physical addr = 12294.  
Virtual addr is 20683: Page# = 5 & Offset = 203. Physical addr = 4299.  
Virtual addr is 18515: Page# = 4 & Offset = 2131. Physical addr = 2131.  
Virtual addr is 28781: Page# = 7 & Offset = 109. Physical addr = 20589.

Virtual addr is 24462: Page# = 5 & Offset = 3982. Physical addr = 8078.  
 Virtual addr is 16399: Page# = 4 & Offset = 15. Physical addr = 15.  
 Virtual addr is 20815: Page# = 5 & Offset = 335. Physical addr = 4431.  
 Virtual addr is 18295: Page# = 4 & Offset = 1911. Physical addr = 1911.  
 Virtual addr is 12218: Page# = 2 & Offset = 4026. Physical addr = 16314.  
 Virtual addr is 13000: Page# = 3 & Offset = 712. Physical addr = 29384.  
 Virtual addr is 12229: Page# = 2 & Offset = 4037. Physical addr = 16325.  
 Virtual addr is 27966: Page# = 6 & Offset = 3390. Physical addr = 11582.  
 Virtual addr is 24894: Page# = 6 & Offset = 318. Physical addr = 8510.  
 Virtual addr is 28929: Page# = 7 & Offset = 257. Physical addr = 20737.  
 Virtual addr is 27865: Page# = 6 & Offset = 3289. Physical addr = 11481.  
 Virtual addr is 5000: Page# = 1 & Offset = 904. Physical addr = 17288.  
 Virtual addr is 2315: Page# = 0 & Offset = 2315. Physical addr = 26891.

### **Notes on Bit wise Operators in C**

1. Bitwise operators work on **bits** and perform bit-by-bit operation.
2. Binary AND Operator (&) - copies a bit to the result if it exists in **both** operands.
3. Binary OR Operator (|) - copies a bit if it exists in **either** operand.
4. Binary Left Shift Operator (<<) - The left operands value is moved left by the number of bits specified by the right operand.
  - a. For example,  $1100 \ll 2 = 110000$
5. Binary Right Shift Operator (>>) - The left operands value is moved right by the number of bits specified by the right operand.
  - a. For example,  $1100 \gg 2 = 11$
6. Assume you have a 8 bit logical address and page size = 16 bytes =  $2^4$ .
7. Number of bits to represent page number = 4
8. Number of bits to represent page offset = OFFSET\_BITS = 4
9. **Page number = logical address >> OFFSET\_BITS**
10.  $\text{OFFSET\_MASK} = 2^4 - 1 = 15$  (in binary it is 00001111)
11. **Page Offset = logical address & OFFSET\_MASK**
12. Suppose the page is stored in the frame "frame\_number", then  
**Physical Address = (frame\_number << OFFSET\_BITS) | offset**

## PART II

In this lab you are going to study Memory Mapping files in C.

You are to write a C program (named labMMF.c) that opens a binary file named **numbers.bin** and maps it to a memory region using the system call **mmap()**. The 'numbers.bin' can be downloaded from Avenue -> Content -> Labs -> Lab 3 material. This file contains **ten** 4-byte integers in binary format.

After memory mapping this file you are to read the contents from this memory mapped region; that is, read one integer at a time, and copy it to an integer array (names **intArray**) using the **memcpy()** function. Finally, you are to loop through the **intArray** array to add all the numbers of the array and output the sum to the console.

In particular, your program needs to do the following:

1. To be able to use **mmap()** system call and the **memcpy()** function you need to add the below header files in addition to your standard input output header file.

```
#include <sys/mman.h> /*For mmap() function*/
#include <string.h> /*For memcpy function*/
```

2. Other useful header files are listed below. They enable you to use the **open()** system call which is used to open a new file and obtain its file descriptor.

```
#include <fcntl.h> /*For file descriptors*/
#include <stdlib.h> /*For file descriptors*/
```

3. Define global variables for the integer array and a signed character pointer to store the starting address of the memory mapped file. E.g.:

```
int intArray[MEMORY_SIZE]; int intArray[INT_COUNT];
signed char *mmapfptr;
```

You can define 'MEMORY\_SIZE' as a macro definition. It is the total number of bytes you will be copying from numbers.bin file. Sample code:

```
#define INT_SIZE 4 // Size of integer in bytes
#define INT_COUNT 10 #define MEMORY_SIZE INT_COUNT * INT_SIZE
```

4. Open the file (**numbers.bin**) using the **open()** system call. Since you will be simply reading this file use the **O\_RDONLY** option. E.g.:

```
int mmapfile_fd = open("numbers.bin", O_RDONLY);
```

5. Use the **mmap()** system call to memory map this file. E.g.:

```
mmapfptr = mmap(0, MEMORY_SIZE, PROT_READ, MAP_PRIVATE,
mmapfile_fd, 0);
```

6. Retrieve the contents of the memory mapped file (using a loop) and store it in the integer array using **memcpy()** function. Sample code to use **memcpy()** is:

```
memcpy(intArray + i, mmapfptr + 4*i, INT_SIZE);
```

**INT\_SIZE** = the size of the contents in bytes to be copied from the memory mapped file to **intArray**. Since we are reading only 4 bytes (size of an integer) at a time, **INT\_SIZE** = 4.

7. Unmap the memory mapped file using the **unmap()** system call. E.g.:

```
munmap(mmapfptr, MEMORY_SIZE);
```

8. Loop through **intArray** to add all numbers in the array and output this sum to the console.

9. Compile your program without errors.

10. Show the program output to your TA.

---

Sample Program output: **./labMMF**  
Sum of numbers = 92